

Machine learning applied to population genomics

Manolo Fernandez Perez

Department of Life Sciences, Imperial College London

@ManoloLearning 

manolofperez@gmail.com 

sites.google.com/site/manolofperez 

Goals

- Conceive and simulate genetic data under competing demographic scenarios
- Understand deep learning background and how a CNN works
- Use CNN to detect regions with selective sweeps on real genomes
- How to use deep learning to compare demographic scenarios



Program

Program

- ***Part I: - The building blocks of a CNN script.***
- ***Practical: Comparing demographic scenarios and detecting selection with deep learning.***
- ***Part II: Quick overview of other applications and future perspectives.***
- ***Part III: Wrapup.***

Part I: The building blocks of a CNN script



The diagram illustrates the building blocks of a CNN script. It starts with an 'Input Patch' of size 28x28. This is followed by a 'Conv' layer with a 5x5 kernel, resulting in 20 feature maps. A 'Max Pool' layer with a 2x2 kernel is applied, resulting in 20 feature maps. This is followed by another 'Conv' layer with a 5x5 kernel, resulting in 50 feature maps. A 'Max Pool' layer with a 2x2 kernel is applied, resulting in 50 feature maps. This is followed by a 'Fully Connected' layer with 500 units. The final output is passed through a 'ReLU' activation function.

CNN Script

Inputs:

Scenario 1




Samples

SNPs

-1	-1	1	0
1	-1	1	1
-1	0	1	-1
-1	1	1	1
1	1	-1	0
0	1	-1	-1

-1	-1	1	0
1	-1	1	1
-1	0	1	-1
-1	1	1	1
1	1	-1	0
0	1	-1	-1


Scenario 2



-1	-1	1	0
1	-1	1	1
-1	0	1	-1
-1	1	1	1
1	1	-1	0
0	1	-1	-1

-1	-1	1	0
1	-1	1	1
-1	0	1	-1
-1	1	1	1
1	1	-1	0
0	1	-1	-1

Scenario 3



-1	-1	1	0
1	-1	1	1
-1	0	1	-1
-1	1	1	1
1	1	-1	0
0	1	-1	-1

-1	-1	1	0
1	-1	1	1
-1	0	1	-1
-1	1	1	1
1	1	-1	0
0	1	-1	-1

Parameters

	Theta	T1	T2	T3	Ne
Sim1					
Sim2					
Sim3					
Sim4					

No of simulations

3-D Numpy array

CNN Script




required python modules.



TensorFlow

Define the CNN architecture.



Load and process the training data.

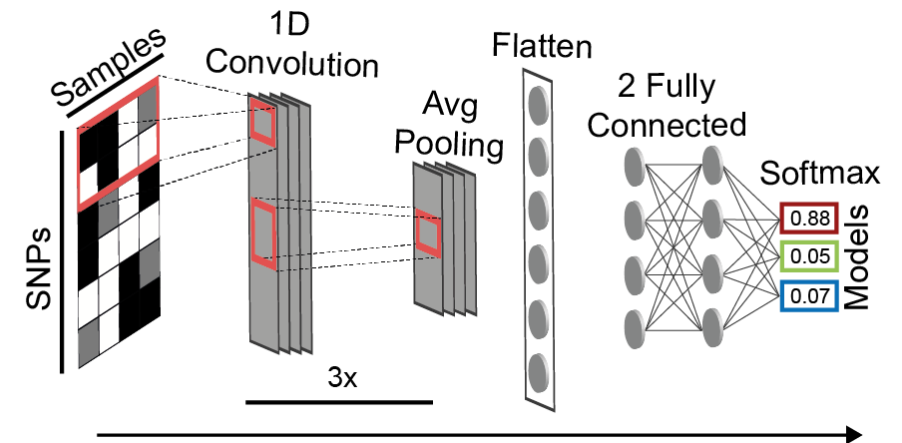
Train the network.

Load the test data and perform cross-validation.

Predict the most likely model for the empirical data

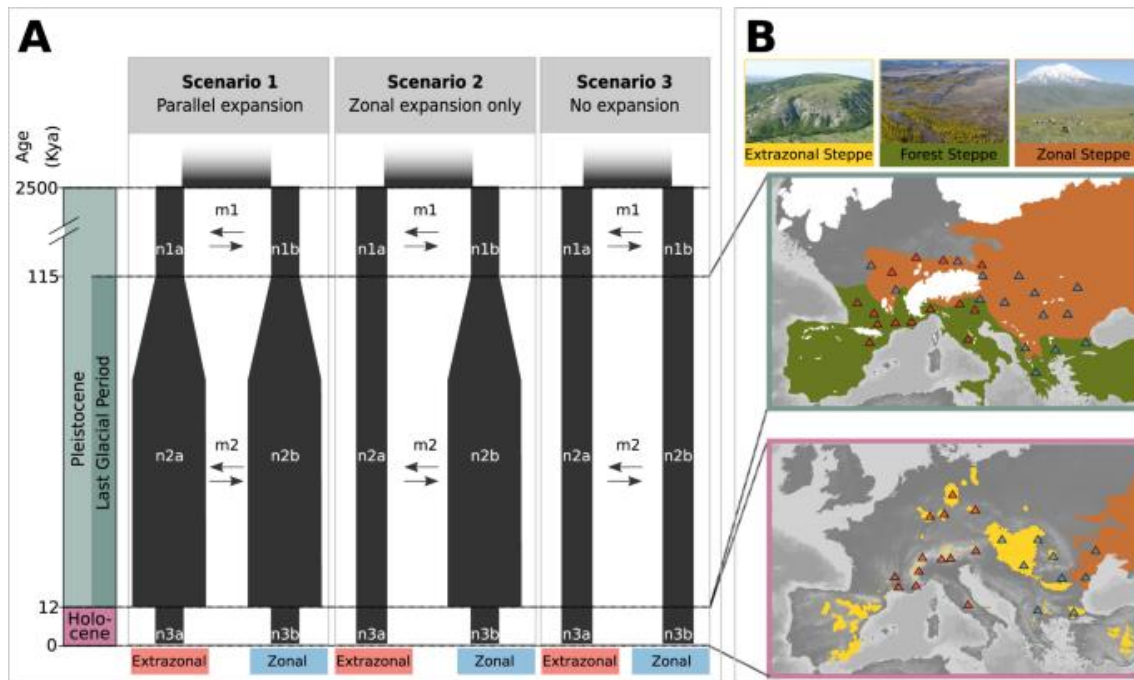
- **Practical Exercise 1:**

- Go through Section 1 of the Part1 script (Demographic models) and try to recognize all the elements of the network. Do you remember the function of each of those elements? Remember that you can add annotations to the code using `#` and add information that might help you when you get back to the script in the future.
- Now run all the cells until you reach the end of section 2. Your network will be training, so now we will have some time to discuss and do a quick review on the CNN elements.

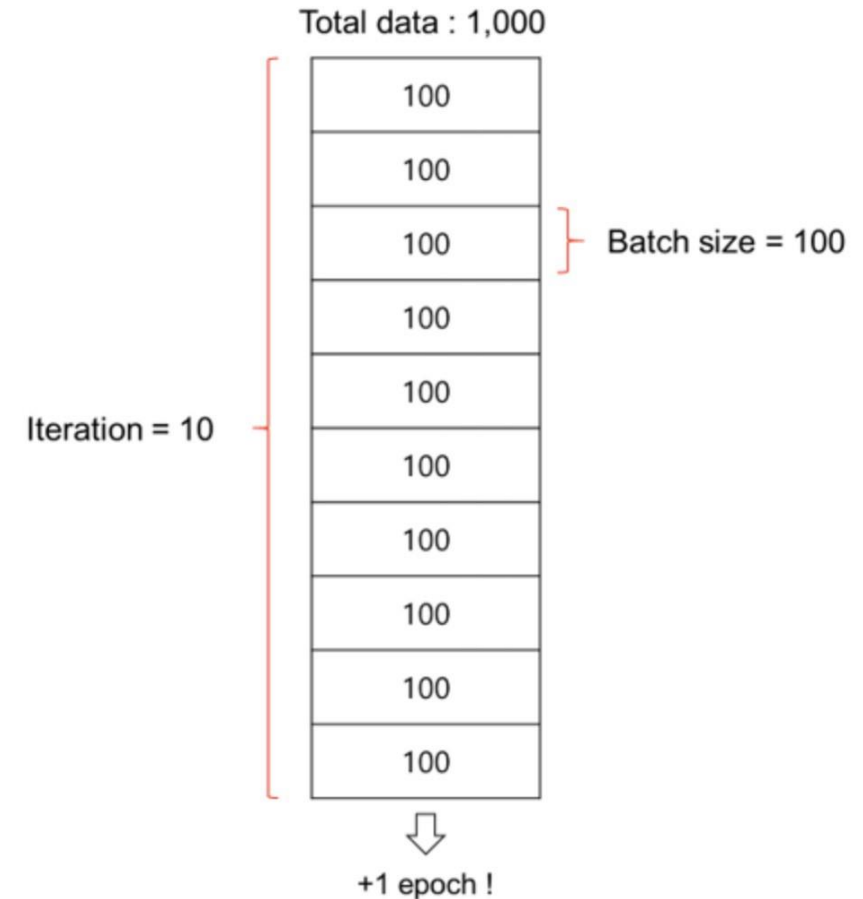



```
# Define parameters for the CNN run.
batch_size = 128
### how much iterations to train the network
epochs = 50
```

```
###n of models
num_classes = 3
```



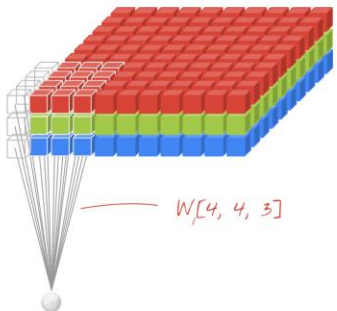
CNN Script



<https://jerryan.medium.com/batch-size-a15958708a6>

CNN Script

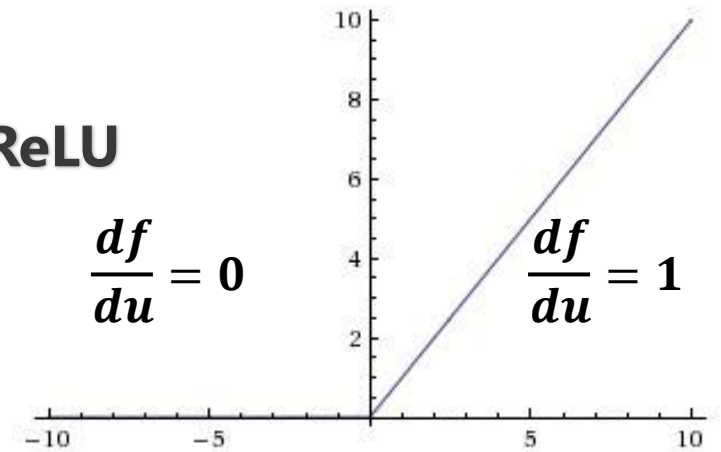
```
# Define the CNN architecture.
def create_cnn(xtest):
    inputShape = (xtest.shape[1], xtest.shape[2])
    ## image size. images need to have EXACTLY the same size
    inputs = Input(shape=inputShape)
    x = inputs
    ## 1D convolution - less computational intensive and is also invariant to the samples order;
    x = Conv1D(256, kernel_size=2, activation='relu', input_shape=(xtest.shape[1], xtest.shape[2]))(x)
    ### Enables the network to learn more complex features / shapes.
    x = AveragePooling1D(pool_size=2)(x)
    x = BatchNormalization()(x)
```



ReLU

$$\frac{df}{du} = 0$$

$$\frac{df}{du} = 1$$



<https://www.kaggle.com/code/dansbecker/rectified-linear-units-relu-in-deep-learning/notebook>

Kirschner et al. (2022) *Nat Comm*

```
# Define the CNN architecture.
```

```
def create_cnn(xtest):
```

```
    inputShape = (xtest.shape[1], xtest.shape[2])
```

```
    ## image size. images need to have EXACTLY the same size
```

```
    inputs = Input(shape=inputShape)
```

```
    x = inputs
```

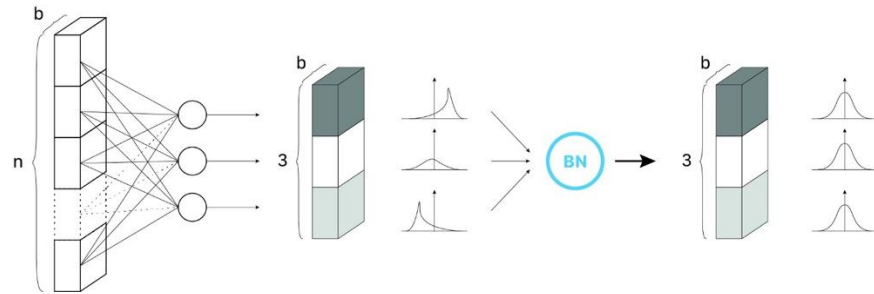
```
    ## 1D convolution - less computational intensive and is also invariant to the samples order;
```

```
    x = Conv1D(256, kernel_size=2, activation='relu', input_shape=(xtest.shape[1], xtest.shape[2]))(x)
```

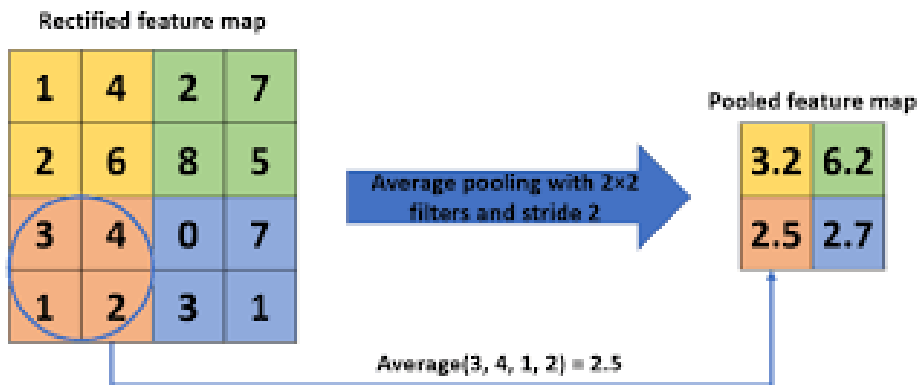
```
    ### Enables the network to learn more complex features / shapes.
```

```
    x = AveragePooling1D(pool_size=2)(x)
```

```
    x = BatchNormalization()(x)
```




<https://towardsdatascience.com/batch-normalization-in-3-levels-of-understanding-14c2da90a338>



Gholamalinezhad & Khosravi (2020) *arXiv*


CNN Script





https://leonardoaraujosantos.gitbook.io/artificial-intelligence/machine_learning/deep_learning/pooling_layer

Kirschner et al. (2022) *Nat Comm*

Linearising the image as in the initial step.



CNN Script



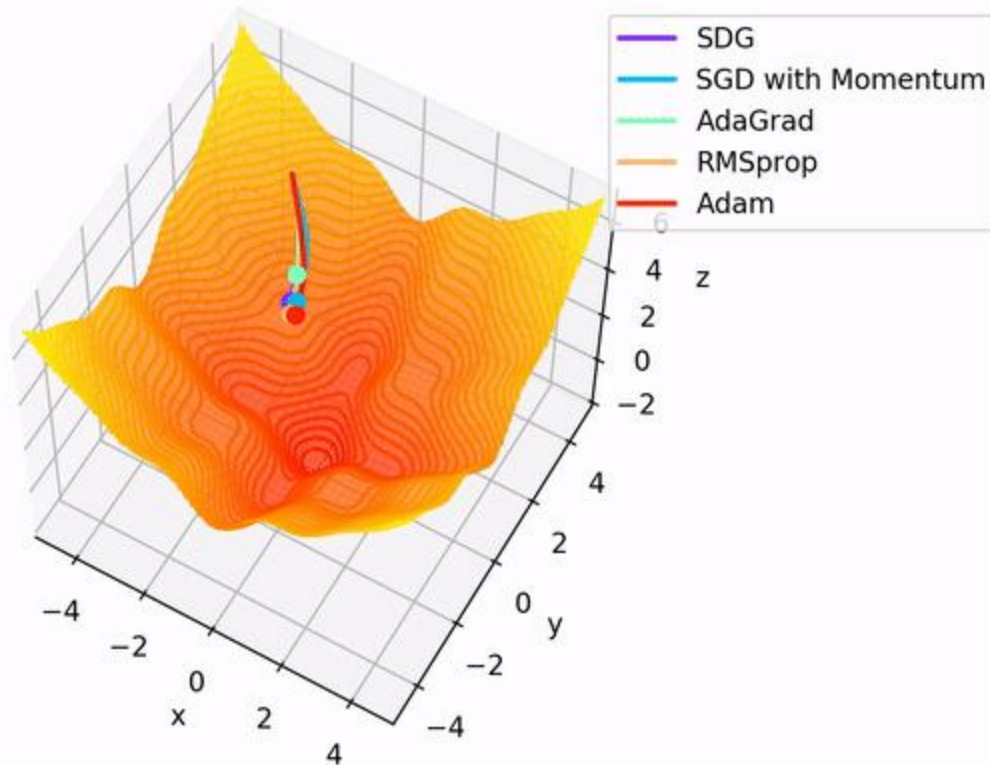
(a) Standard Neural Net

CNN Script


```
# Compile the CNN.
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer='Adam',
              metrics=['accuracy'])

# We will use early stopping and save the model with the best val_accuracy.
earlyStopping = EarlyStopping(monitor='val_accuracy', patience=25, verbose=0, mode='max', restore_best_weights)
### stop training when validation error increases (wait 25 epochs to see if there is any improvement).
```

Optimizer Comparison



<https://towardsdatascience.com/complete-guide-to-adam-optimization-1e5f29532c3d>





<https://towardsdatascience.com/a-practical-introduction-to-early-stopping-in-machine-learning-550ac88bc8fd>

#Run the CNN

```
history = model.fit(xtrain, ytrain, batch_size=batch_size,  
                    epochs=epochs,  
                    verbose=1,  
                    validation_data=(xval, yval), callbacks=[earlyStopping])
```

CNN Script




<https://jerryan.medium.com/batch-size-a15958708a6>

- Practical Exercise 2:

Now open the Part 2 script. It uses Matteo's software (ImaGene) that is a CNN approach to infer selection at the LCT locus.

It uses a "simulation-on-the-fly"-like approach: training data is continuously generated by simulations to avoid the network to see the same data twice and therefore to reduce overfitting. This is a valuable consideration since, when reliable simulators are available, we have access to theoretically infinite training data, the latter being constrained by computing time only.

There are also functions to automate some of the steps we did in the previous example. You can compare the strategies and the architectures adopted.



The background is a dark, textured surface with various colorful, glowing lines and shapes. There are several long, thin lines in shades of blue, green, and yellow, some of which are curved or bent. There are also some rectangular shapes and clusters of small, glowing points. The overall effect is a futuristic or digital aesthetic.

Part II: Quick overview of other applications and future perspectives.

Deep Learning in Population Genetics

Kevin Korfmann¹, Oscar E. Gaggiotti², and Matteo Fumagalli ^{3,*}

¹Professorship for Population Genetics, Department of Life Science Systems, Technical University of Munich, Germany

²Centre for Biological Diversity, Sir Harold Mitchell Building, University of St Andrews, Fife KY16 9TF, UK

³Department of Biological and Behavioural Sciences, Queen Mary University of London, UK

*Corresponding author: E-mail: m.fumagalli@qmul.ac.uk.

Accepted: 16 January 2023

The Unreasonable Effectiveness of Convolutional Neural Networks in Population Genetic Inference

Lex Flagel,^{1,2} Yaniv Brandvain,² and Daniel R. Schrider^{*,3}

¹Monsanto Company, Chesterfield, MO

²Department of Plant and Microbial Biology, University of Minnesota, St. Paul, MN

³Department of Genetics, University of North Carolina, Chapel Hill, NC

*Corresponding author: E-mail: drs@unc.edu.

Associate editor: Yuseob Kim

Review Article | Published: 04 September 2023

Harnessing deep learning for population genetic inference

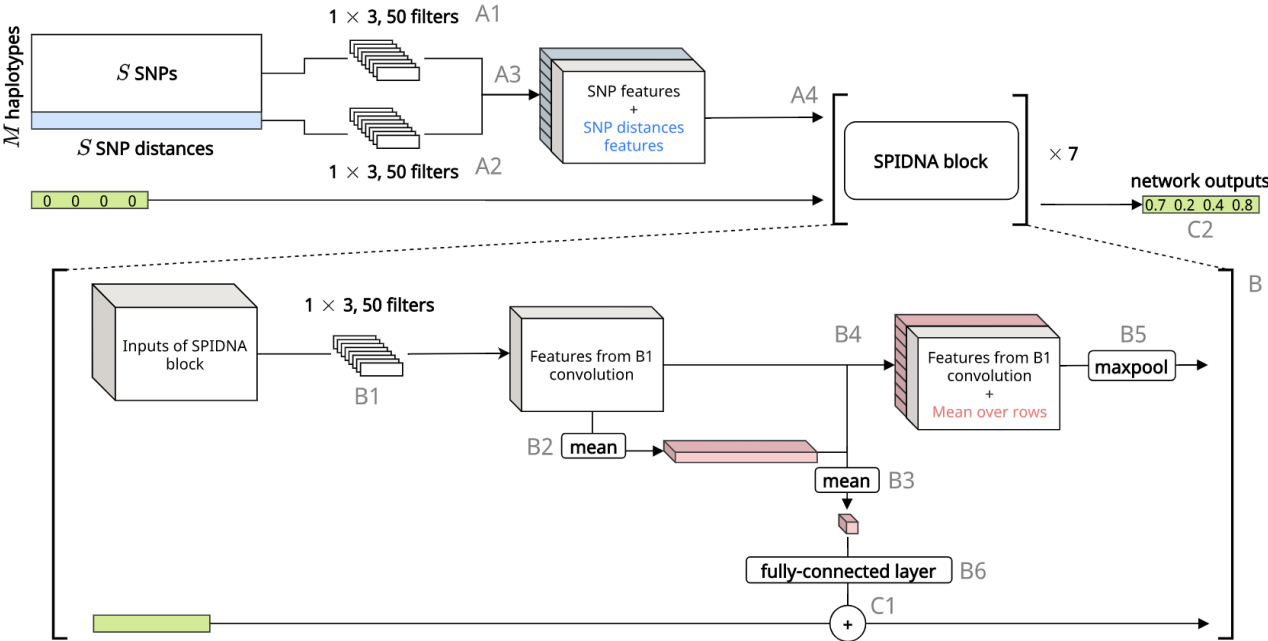
[Xin Huang](#) , [Aigerim Rymbekova](#), [Olga Dolgova](#), [Oscar Lao](#)  & [Martin Kuhlwilm](#) 

[Nature Reviews Genetics](#) **25**, 61–78 (2024) | [Cite this article](#)

8148 Accesses | **4** Citations | **41** Altmetric | [Metrics](#)

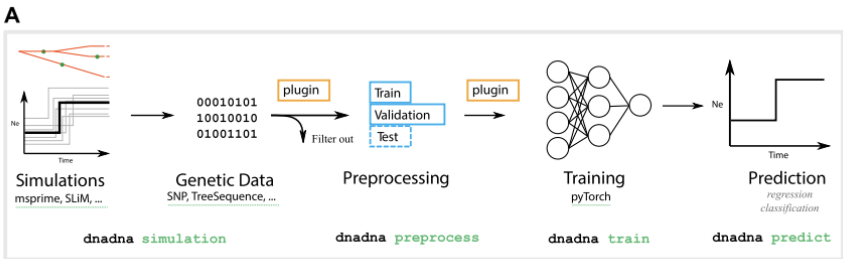
Deep learning for population size history inference: Design, comparison and combination with approximate Bayesian computation

Théophile Sanchez  | Jean Cury  | Guillaume Charpiat | Flora Jay 



Genetics and population analysis **dnadna: a deep learning framework for population genetics inference**

Théophile Sanchez^{1†}, Erik Madison Bray^{1†}, Pierre Jobic^{1,2}, Jérémy Guez^{1,3}, Anne-Catherine Letournel¹, Guillaume Charpiat¹, Jean Cury ^{1,4*‡} and Flora Jay ^{1*‡}



B Standard workflow: train a newly implemented network on existing simulations

```
1/ dnadna preprocess Demo_preprocessing_config.yaml --dataset-config Demo_dataset_config.yaml
2/ dnadna train Demo_training_config.yaml --plugin local_net.py
3/ dnadna predict run_xxx/Demo_run_xxx_best_net.pth Testset/*/*.npz --plugin local_net.py
```

describes a previously simulated training set
- Try new architecture
- Update hyperparameters

D Example of a training config file

```
# the simulation configuration
simulation:
  id: demo_simulation_config.yaml
  learned_params:
    event_time:
      type: regression
      log_transform: true
      log_prior: MSE
    event_size:
      type: regression
      log_transform: true
      log_prior: MSE
  networks:
    name: myNet
    param: 3
    n_epochs: 5
    optimizer:
      name: Adam
      param:
        learning_rate: 0.001
        weight_decay: 0.1
```

C Standard workflow: reuse a trained network on one's dataset

```
1/ dnadna predict pretrained_SPIDNA_net.pth myData/*.npz --preprocessing
```


Contains optimized weight, and all config parameters used for training. Contains means and std to unstandardize prediction
Apply same preprocessing e.g. filter out sequences with less than X SNPs and N individuals

E Example of a network plugin

```
from dnadna import nets
from torch.nn.functional import relu

class myNet(nets.Network):
    def __init__(self, param):
        super().__init__(param)
        self.param = param
    def forward(self, x):
        ...
```

Only change compared to using only pytorch



G Classification task



```
dnadna train classif_config.yaml
dnadna predict classif_best_net.pth test_set/*.npz
```

	Expected	Expansion	Decline
Observed			
Expansion		970	194
Decline		100	736

Accuracy = 85.3%


Automatic inference of demographic parameters using generative adversarial networks

Zhanpeng Wang¹ | Jiaping Wang¹ | Michael Kourakos² | Nhung Hoang² |
Hyong Hark Lee² | Iain Mathieson³ | Sara Mathieson¹ 


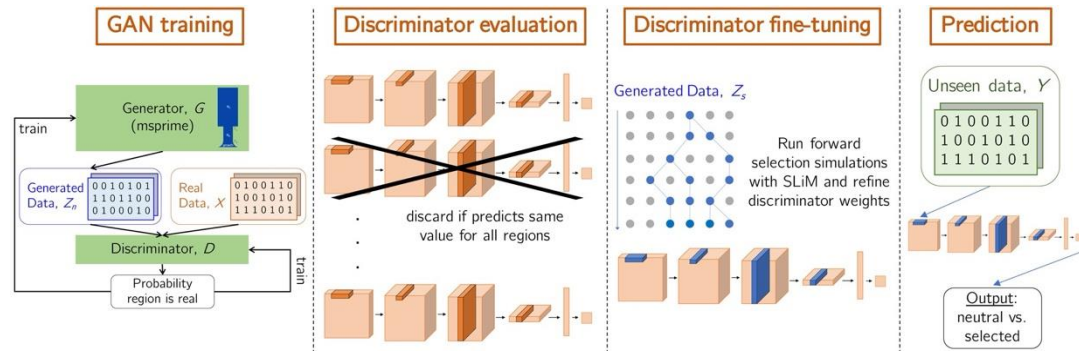


Automatic inference of demographic parameters using generative adversarial networks

Zhanpeng Wang¹ | Jiaping Wang¹ | Michael Kourakos² | Nhung Hoang² |
Hyong Hark Lee² | Iain Mathieson³ | Sara Mathieson¹ 



Neural Network for Genomic data



Peer Community Journal

Section: Evolutionary Biology

Research article

Published
2024-03-18

Cite as

Kevin Korfmann, Thibaut Paul Patrick Sellinger, Fabian Freund, Matteo Fumagalli and Aurélien Tellier (2024) *Simultaneous Inference of Past Demography and Selection from the Ancestral Recombination Graph under the Beta Coalescent*, Peer Community Journal, 4: e33.

Simultaneous Inference of Past Demography and Selection from the Ancestral Recombination Graph under the Beta Coalescent

Kevin Korfmann ,^{#,1}, Thibaut Paul Patrick Sellinger ,^{#,2,1}, Fabian Freund ,^{3,4}, Matteo Fumagalli ,^{5,6}, and Aurélien Tellier ,¹


5. Visualization of inferred variables

4. Masking of time-relevant regions and column-wise mean


3. Last pooling step with feature vector containing inferred variables

2. Learned subgraph with updated feature vectors

1. Coalescent trees with feature vectors



Integrative Deep Learning species delimitation




Simulate SNPs
and tree

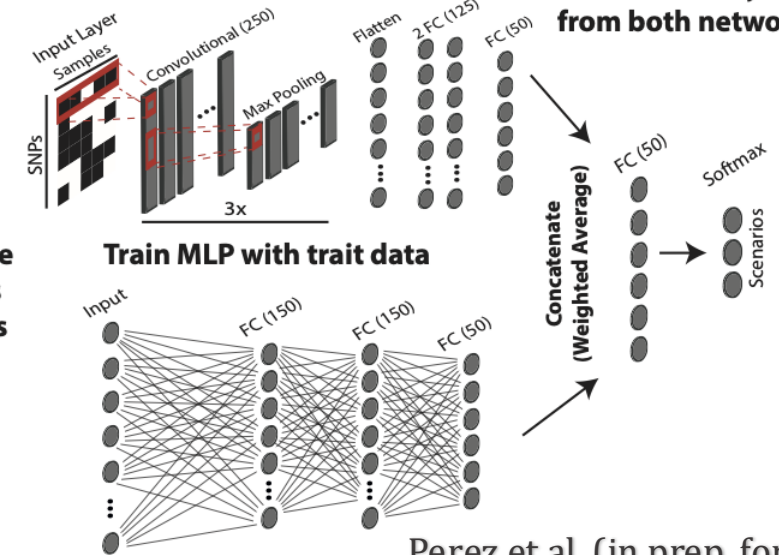
Transform SNPs
to image

Train CNN with SNP data

Combine the
Dense (FC) layers
from both networks



Simulate discrete
and continuous
traits from trees




Perez et al. (in prep. for *Syst. Biol.*)


Deep Learning for **phylogenetics** and **macroevolution**

Perez & Gascuel (in prep. for *Syst. Biol.*)


BD



BDEI


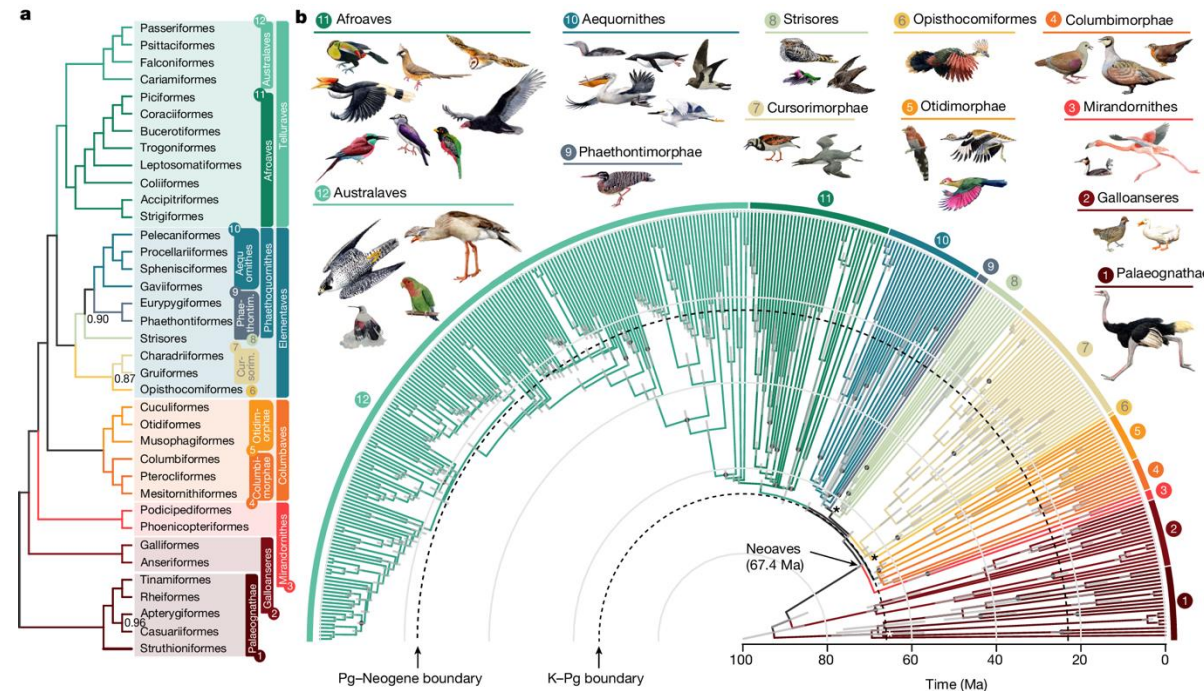


BDSS



Voznica et al. (2022) Nat Comm

Stiller et al 2024 Nature



$$q = q_{01} = q_{10}$$

$$\varepsilon = \mu_0 / \lambda_0 = \mu_1 / \lambda_1$$

What's next?

Your
Project
Here



Population Genomics: background and tools

18 – 26 May 2017 | Napoli, Italy

Part III: Wrapup.

Goals

- Conceive and simulate genetic data under competing demographic scenarios
- Understand deep learning background and how a CNN works
- Use CNN to detect regions with selective sweeps on real genomes
- How to use deep learning to compare demographic scenarios

